

STP 598: Homework 3

Antonio Campbell, Sinta Sulistyo, Atta Ullah, Penny Wu

2/22/2021

1. Using Cross Validation

We are going to use the used cars data again. Previously, we used the “eye-ball” method to choose k for a kNN fit for mileage predicting price. Use 5-fold cross-validation to choose k . How does your fit compare with the eyeball method?

| price | mileage | year |
|-------|---------|------|
| 43995 | 36858 | 2008 |
| 44995 | 46883 | 2012 |
| 25999 | 108759 | 2007 |
| 33880 | 35187 | 2007 |
| 34895 | 48153 | 2007 |
| 5995 | 121748 | 2002 |


Plot the data and then add the fit using the k you chose using cross-validation and the k you choose by eye-ball.

```
#get variables we want
x1 = cbind(carDat$mileage)
colnames(x1) = c("mileage")
y = carDat$price

#run cross val several times
set.seed(99)
kv = 2:50 #these are the k values (k as in kNN) we will try
cvmean = rep(0,length(kv)) #will keep average rmse here
ndocv = 5 #number of CV splits to try
n=length(y)
cvmat = matrix(0,length(kv),ndocv) #keep results for each split


for(i in 1:ndocv) {
  cvtemp = docvknn(x1,y,kv,nfold=5)
  cvmean = cvmean + cvtemp
  cvmat[,i] = sqrt(cvtemp/n)
}

cvmean = cvmean/ndocv
cvmean = sqrt(cvmean/n)
```



```
kbest2 = kbest = kv[which.min(cvmean)]
cat("the best k is: ",kbest,"\n")
```

```
## the best k is: 20
```



Comparing our eyeball fit to the one provided by cross validation we see the two are pretty similar, but we favor the model that is more complicated with $k = 20$. The two fit RMSEs follow:

| EyeballedK | BestK |
|------------|----------|
| 9013.226 | 8810.022 |

We see that as the name suggests, we have the best fit if we use the k chosen through cross validation.

Use kNN with the k you chose using cross-validation to get a prediction for a used car with 100,000 miles on it. Use all the observations as training data to get your prediction (given your choice of k).



```
near3 = knn(y~x1,data.frame(x1, y),data.frame(x1=100000),k=20,kernel = "rectangular")
cat("knn predicted value: ",near3$fitted,"\n")
```

```
## knn predicted value: 18251.55
```

2. kNN, Cars Data with Mileage and Year

Use kNN to get a prediction for a 2008 car with 75,000 miles on it! Remember to use cross-validation to choose k and scale your x 's !!

```
#get variables we want
x2 = cbind(carDat$mileage, carDat$year)
colnames(x2) = c("mileage","year")
mmssc=function(x) {return((x-min(x))/(max(x)-min(x)))}
xs = apply(x2,2,mmssc) #apply scaling function to each column of x
```



```
#run cross val several times
kv = 2:35 #these are the k values (k as in kNN) we will try
cvmean = rep(0,length(kv)) #will keep average rmse here
ndocv = 5 #number of CV splits to try
n=length(y)
cvmat = matrix(0,length(kv),ndocv) #keep results for each split
```


```

for(i in 1:ndocv) {
  cvtemp = docvknn(xs,y,kv,nfold=5)
  cvmean = cvmean + cvtemp
  cvmat[,i] = sqrt(cvtemp/n)
}
cvmean = cvmean/ndocv
cvmean = sqrt(cvmean/n)

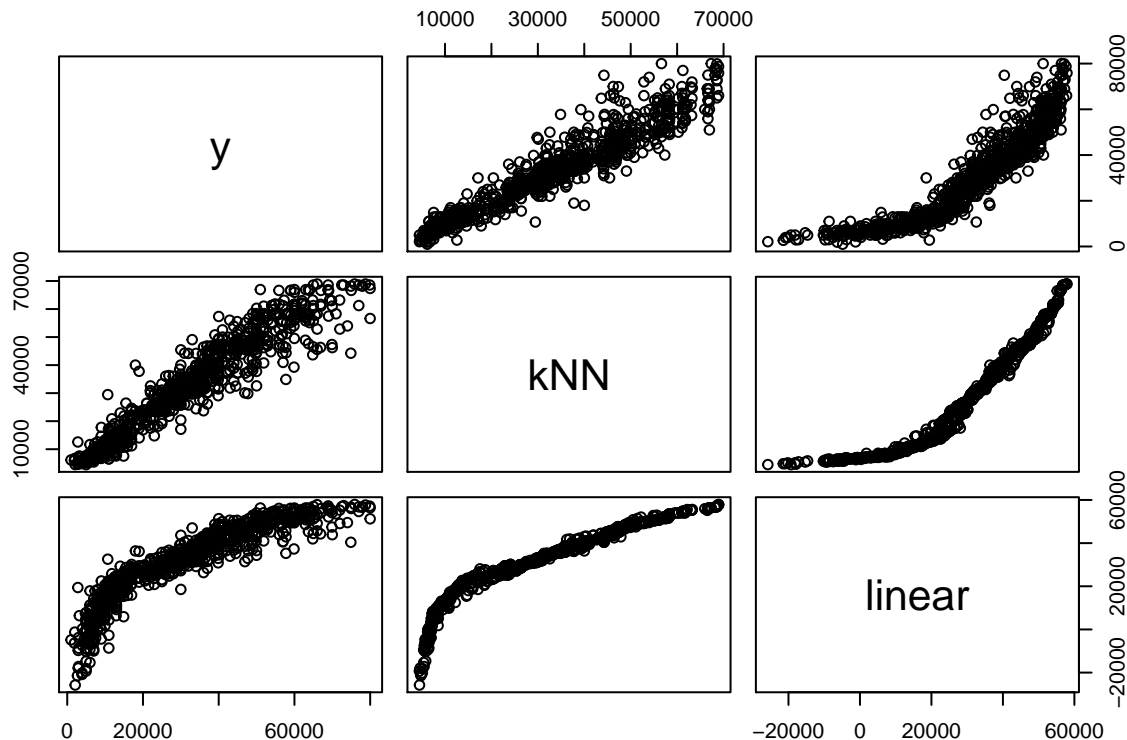
plot(log(1/kv),cvmean,type="n",ylim=range(cvmat),xlab="k",cex.lab=1.5)

for(i in 1:ndocv) lines(log(1/kv),cvmat[,i],col=i,lty=3) #plot each result
lines(log(1/kv),cvmean,type="b",col="black",lwd=2) #plot average result

```



```
## the best k is: 26
```



```
##           y      kNN    linear
## y      1.000000 0.9581329 0.9123897
## kNN    0.9581329 1.0000000 0.9525080
## linear 0.9123897 0.9525080 1.0000000
```

```
## knn predicted value: 31646.65
```

Is your predictive accuracy better using (mileage,year) than it was with just mileage?

Comparing the RMSEs we have:

```
kbest2 = kkn(y~x1,data.frame(x1, y),data.frame(x1), k=kbest2, kernel = "rectangular")
k = rmse(kbest2$fitted.values,y)
j = rmse(near2$fitted.values,y)
kable(cbind(Mileage = k, MileageYear = j))
```

| Mileage | MileageYear |
|----------|-------------|
| 8810.022 | 5275.12 |

So yes, the predictive accuracy is quite better.

3. Choice of Kernel

In our notes examples we used `kernel="rectangular"` when calling the R function `kkn`. So, you can weight the `y` values at the neighbors equally, or weight the closer ones more heavily. Typically default is equal weights. Using the used cars data and predictors (features!!) (mileage,year) try a weighting option other than uniform.

We try some of the different weighting methods and report their RMSEs below.

| Triangular | Triweight | Optimal | Normal |
|------------|-----------|----------|---------|
| 5062.027 | 4838.124 | 5062.027 | 5275.12 |

In this case, we see that tri-weight has the best accuracy among these options.